

Network Design with Coverage Costs

Siddharth Barman ¹ Shuchi Chawla ² Seeun William Umboh ²

¹Caltech

²University of Wisconsin-Madison

APPROX-RANDOM 2014


Motivation

Physical Flow vs Data Flow



Commodity Network


vs.



Internet

Unlike physical commodities, data can be easily duplicated, compressed, and combined.

Physical Flow vs Data Flow



Unlike physical commodities, data can be easily duplicated, compressed, and combined.

Why Coverage Costs?


- Want a cost structure that captures bandwidth savings obtained by eliminating **redundancy in data**.
 - Packet deduplication capability exists in networks today.

Problem Statement

Load Function


We focus on redundant-data elimination:

Load on an edge $\ell_e = \#$ **distinct** data packets on it.





Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$





Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$
- A global set of packets Π .



Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$
- A global set of packets Π .
- A demand set $D_j \subset \Pi$ for each j .





Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$
- A global set of packets Π .
- A demand set $D_j \subset \Pi$ for each j .

Goal: Find (s_j, t_j) path on which to route D_j for each j minimizing

$$\sum_{\text{edges}} \text{Cost of edge} \times \text{Load on edge}$$





Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$
- A global set of packets Π .
- A demand set $D_j \subset \Pi$ for each j .

Goal: Find (s_j, t_j) path on which to route D_j for each j minimizing

$$\sum_{\text{edges}} \text{Cost of edge} \times \# \text{ distinct packets on edge}$$





Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$
- A global set of packets Π .
- A demand set $D_j \subset \Pi$ for each j .

Goal: Find (s_j, t_j) path on which to route D_j for each j minimizing

$$\sum_{\text{edges}} \text{Cost of edge} \times \# \text{ distinct packets on edge}$$





Input:

- Graph with edge costs, g terminal pairs $(s_1, t_1), \dots, (s_g, t_g)$
- A global set of packets Π .
- A demand set $D_j \subset \Pi$ for each j .

Goal: Find (s_j, t_j) path on which to route D_j for each j minimizing

$$\sum_{\text{edges}} \text{Cost of edge} \times \# \text{ distinct packets on edge}$$



More generally, we consider terminal *groups*.

Input:



- Graph with edge costs, g terminal groups $X_1, \dots, X_g \subset V$.
- A global set of packets Π .
- A demand set $D_j \subset \Pi$ for each j .

Goal: Find a Steiner tree for X_j on which to broadcast D_j for each j minimizing

$$\sum_{\text{edges}} \text{Cost of edge} \times \# \text{ distinct packets on edge}$$



Special cases (for terminal pairs):

- Pairwise disjoint demands \implies shortest-paths





Special cases (for terminal pairs):

- Pairwise disjoint demands \implies shortest-paths




- Identical demands \implies Steiner forest



Challenges

- Intuitively, difficulty depends on complexity of demand sets


Π



Challenges

- Intuitively, difficulty depends on complexity of demand sets

Π




- Desire an approximation in terms of g (number of groups), e.g. $O(\log g)$ (tree embeddings only yield $O(\log n)$).

Our results


Laminar Demands

The family of demand sets is said to be **laminar** if $\forall i, j$ one of the following holds: $D_i \cap D_j = \emptyset$ or $D_i \subset D_j$ or $D_j \subset D_i$.



Laminar Demands

The family of demand sets is said to be **laminar** if $\forall i, j$ one of the following holds: $D_i \cap D_j = \emptyset$ or $D_i \subset D_j$ or $D_j \subset D_i$.




Theorem

NDCC with laminar demands admits a 2-approximation.

Previous work: $O(\log |\Pi|)$ [Barman-Chawla '12].


Sunflower Demands

The family of demand sets is said to be a **sunflower** family if there exists a core $C \subset \Pi$ such that $\forall i, j, D_i \cap D_j = C$.



Sunflower Demands


The family of demand sets is said to be a **sunflower** family if there exists a core $C \subset \Pi$ such that $\forall i, j, D_i \cap D_j = C$.



No $O(\log^{1/4-\gamma} g)$ -approximation for any $\gamma > 0$ (reduction from single-cable buy-at-bulk [Andrews-Zhang '02]).

Sunflower Demands

The family of demand sets is said to be a **sunflower** family if there exists a core $C \subset \Pi$ such that $\forall i, j, D_i \cap D_j = C$.



No $O(\log^{1/4-\gamma} g)$ -approximation for any $\gamma > 0$ (reduction from single-cable buy-at-bulk [Andrews-Zhang '02]).

Theorem

NDCC with sunflower demands admits an $O(\log g)$ approximation over unweighted graphs with $V = \bigcup_j X_j$.

Related Work

Network design with economies of scale.

$$\sum_{\text{edges}} \text{Cost of edge} \times \text{Load on edge}$$

Related Work

Network design with economies of scale.

$$\sum_{\text{edges}} \text{Cost of edge} \times \text{Load on edge}$$

- Submodular costs on edges [[Hayrapetyan-Swamy-Tardos '05](#)]
 - $O(\log n)$ via tree embeddings.

Related Work


Network design with economies of scale.

$$\sum_{\text{edges}} \text{Cost of edge} \times \text{Load on edge}$$

- Submodular costs on edges [[Hayrapetyan-Swamy-Tardos '05](#)]
 - ▶ $O(\log n)$ via tree embeddings.
- Buy-at-Bulk Network Design [[Salman et al. '97](#)]
 - ▶ Single-source: $O(1)$ [[Guha et al. '01](#), [Talwar '02](#), [Gupta-Kumar-Roughgarden '03](#)]
 - ▶ Multiple-source: $\Omega(\log^{\frac{1}{4}} n)$ hardness [[Andrews '04](#)]

Approach: Laminar

Laminar




- Naive LP does not have much structure
- Key idea: exploit laminarity to write a different LP
- Run an extension of AKR-GW primal-dual algorithm to get 2-approx

**Approach: Sunflower
(single-source for this talk)**

Sunflower Family of Demands

The family of demand sets is said to be a **sunflower** family if there exists $C \subset \Pi$ such that $\forall i, j, D_i \cap D_j = C$.



Theorem

Network Design with Coverage Costs in the sunflower demands setting admits an $O(\log g)$ approximation over unweighted graphs with $V = \bigcup_j X_j$.

Structure of OPT

Define:

- $E_j^* =$ Steiner tree for X_j on which OPT broadcasts $D_j \supset C$
- $E_0^* = \bigcup_j E_j^*$

Structure of OPT

Define:

- $E_j^* =$ Steiner tree for X_j on which OPT broadcasts $D_j \supset C$
- $E_0^* = \bigcup_j E_j^*$

Note: E_0^* spans V because $V = \bigcup_j X_j$ and single-source (each X_j contains s)

Structure of OPT

Define:

- $E_j^* =$ Steiner tree for X_j on which OPT broadcasts $D_j \supset C$
- $E_0^* = \bigcup_j E_j^*$

Note: E_0^* spans V because $V = \bigcup_j X_j$ and single-source (each X_j contains s)

$$\begin{aligned} c(\text{OPT}) &= |C| \cdot c(E_0^*) + \sum_j |D_j \setminus C| \cdot c(E_j^*) \\ &\geq |C| \cdot c(\text{MST}) + \sum_j |D_j \setminus C| \cdot c(\text{opt Steiner for } X_j) \end{aligned}$$

Group Spanners

Definition

For a graph G and g groups $X_1, \dots, X_g \subset V$, a subgraph H is a (α, β) group spanner if

- $c(H) \leq \alpha c(\text{MST})$,
- $c(\text{opt Steiner for } X_j \text{ in } H) \leq \beta c(\text{opt Steiner for } X_j \text{ in } G)$ for all j .

Group Spanners

Definition

For a graph G and g groups $X_1, \dots, X_g \subset V$, a subgraph H is a (α, β) group spanner if

- $c(H) \leq \alpha c(\text{MST})$,
- $c(\text{opt Steiner for } X_j \text{ in } H) \leq \beta c(\text{opt Steiner for } X_j \text{ in } G)$ for all j .

Usual spanners: X_j s are all vertex pairs.

- $c(\text{shortest } u - v \text{ path in } H) \leq \beta c(\text{shortest } u - v \text{ path in } G)$ for all $u, v \in V$.

Using Group Spanners

Given (α, β) group spanner H , route D_j along 2-approximate Steiner tree for X_j in H .

$$\begin{aligned}\text{Cost} &\leq |C| \cdot c(H) + \sum_j |D_j \setminus C| \cdot 2c(\text{opt Steiner for } X_j \text{ in } H) \\ &\leq |C| \cdot \alpha c(\text{MST}) + \sum_j |D_j \setminus C| \cdot 2\beta c(\text{opt Steiner for } X_j) \\ &\leq \max\{\alpha, 2\beta\} \text{OPT}.\end{aligned}$$

Group Spanners Result

Lemma

Given an *unweighted* graph G and g groups $X_1, \dots, X_g \subset V$ with $V = \bigcup_j X_j$, we can construct in polynomial time a $(O(1), O(\log g))$ group spanner.




Theorem

Network Design with Coverage Costs in the sunflower demands setting admits an $O(\log g)$ approximation over unweighted graphs with $V = \bigcup_j X_j$.


Summary

- Coverage cost model for designing information networks.




Summary

- Coverage cost model for designing information networks.




- Laminar demands: 2-approximation via primal-dual.




Summary

- **Coverage cost model** for designing information networks.



- **Laminar demands**: 2-approximation via primal-dual.



- **Sunflower demands**: $O(\log g)$ -approximation (unweighted graph, no Steiner vertices).



- **Group spanners**: $(O(1), O(\log g))$ group spanners (unweighted graph, no Steiner vertices).

Open Problems

- Group spanners: $(O(\log g), O(\log g))$ in general?

Open Problems

- Group spanners: $(O(\log g), O(\log g))$ in general?
- Sunflower demands: $O(\log g)$ approximation in general?

Open Problems

- Group spanners: $(O(\log g), O(\log g))$ in general?
- Sunflower demands: $O(\log g)$ approximation in general?
- Terminal pairs with single-source: $O(1)$?

Open Problems

- Group spanners: $(O(\log g), O(\log g))$ in general?
- Sunflower demands: $O(\log g)$ approximation in general?
- Terminal pairs with single-source: $O(1)$?

Thanks!